

Project: Virtual Memory

**Originally created by
Lubomir Bic**

**Presented/modified by
Michael Dillencourt**

University of California, Irvine

Project Overview

- Implement a **virtual memory** system using **segmentation** and **paging**
 - manage segment and page tables in a simulated physical memory
 - accept virtual addresses and translate them into physical addresses
- Simple version:
 - assumes that entire VA space is resident in physical memory (PM)
 - earns partial credit
- Extended version:
 - supports **demand paging**
 - earns full credit
 - implements **translation look-aside buffer** to improve efficiency
 - not required for this project

Principles of VM

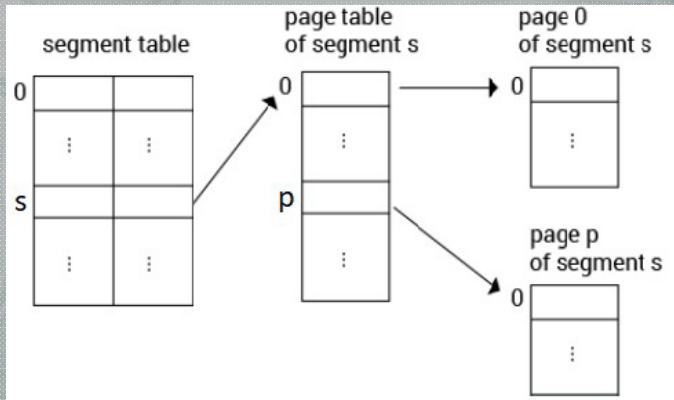
- **VM:** logical address space whose size may exceed size of physical space
 - implementation: segmentation, paging, or both
- **Segmentation**
 - logical address space is divided into variable-size blocks (segments)
 - each segment is contiguous in PM
 - segment table (ST) keeps track of starting addresses of segments
 - $VA = (s, w)$, where s is segment number and w is offset within segment
- **Paging**
 - logical address space is divided into fixed-size blocks (pages)
 - PM is divided into fixed-size blocks (page frames)
 - page table PT keeps track of which page is in which frame
 - $VA = (p, w)$, where p is the page number and w is offset within page

Principles of VM

- Advantage of **segmentation**:
 - each segment corresponds to **logical** component (code, data, stack, etc.)
 - easier for sharing and linking
- Advantage of **paging**:
 - frame size = page size, thus any page may be mapped into any available frame
 - no need to search for and maintain **variable-size** memory partitions
- To gain advantages of both: **combine segmentation with paging**

Segmentation with Paging

- each segment is divided into fixed-size pages
- each ST entry points to PT corresponding to one segment
- each PT entry points to one page

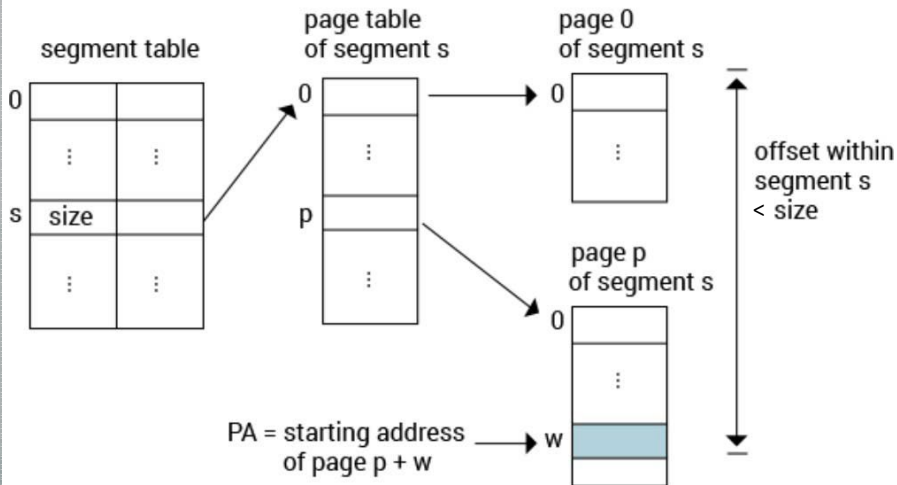


Address Translation

- VA is a nonnegative integer
- Number of bits to represent VA determines size of the VA space
 - Ex: with VA of 32 bits, 2^{32} addresses can be created
- VA = (s, p, w), s is segment number, p is page number, w is offset within page
- Number of bits used to represent s, p, w determine sizes of ST, PT, and page
- PA is also a nonnegative integer
 - Number of bits determines size of PM
- VM manager translates VAs into PAs
 - first, break VA into 3 integer components, s, p, and w
 - use s, p, w as indices into ST, PT, page

Address Translation

$VA = (s, p, w)$



- segments have different sizes
- VA must not exceed size

Demand Paging

- If size of VM exceeds size of PM, then
 - not all pages can be present in PM
 - must be loaded from a disk as needed: **demand paging**
- Demand paging applies also to PTs: PT may not be resident
- **Present bit**
 - associated with each entry of the PT and ST
 - if true, then entry contains **frame** number of the page or PT
 - if false, then entry contains location of page or PT on **disk**
- **Page fault:** reference to nonresident page or PT
 - locate missing page/PT on disk, allocate page frame, load page/PT

Two versions of project

1. Without demand paging (max score is 60%)
2. With demand paging (max score is 100%)

Notes:

- You can submit both versions, and we will use the version that provides the most credit
- This is recommended if you are not confident that your implementation of the demand paging project is reliable.




Basic Version – Without Demand Paging

VM Specs w/o Demand Paging

- Memory is word-addressable, where each word is an integer
- There is only a **single** process and hence only a single ST
- Each ST entry consists of 2 integers:
 - **size** of the segment s (number of words)
 - frame number holding PT of segment s
- If segment does not exist, then both fields in the entry are 0
- Each PT entry contains frame number of page
- VA: 32-bit integer, divided into s , p , w
- Each component occupies 9 bits:
 - PT size = page size = 512 words
 - ST size = 1024 words (each entry occupies 2 integers)

Representation of PM

- PM: array of 524,288 integers, divided into 1024 frames of 512 words each




- PM[2s] refers to size of s, PM[2s+1] refers to frame number of PT
- Given s, p, w: $PA = PM[PM[2s+1]*512+p]*512+w$

Initialization of the PM

- Pre-initialization: $PM[i] = 0$ for all i
- PM is then initialized from a file, which specifies:
 - Line 1: $s_1 z_1 f_1 s_2 z_2 f_2 \dots s_n z_n f_n$ \leftarrow defines ST
 - $s_i z_i f_i$ means: PT of segment s_i resides in frame f_i , length of segment s_i is z_i
 - Ex: 6 3000 4: PT of segment 6 resides in frame 4, size of segment 6 is 3000
Initialize: $PM[2*6] = PM[12] = 3000$ and $PM[2*6+1] = PM[13] = 4$
 - Line 2: $s_1 p_1 f_1 s_2 p_2 f_2 \dots s_m p_m f_m$ \leftarrow defines PTs
 - $s_j p_j f_j$ means: page p_j of segment s_j resides in frame f_j
 - Ex: 6 5 9: page 5 of segment 6 resides in frame 9
Initialize:
 $PM[PM[13]*512+5] = PM[4*512+5] = PM[2053] = 9$

Initialization of the PM

- Line 1: 6 3000 4 (VAs 0-2999 are valid)
- Line 2: 6 5 9 ... (assume pages 0-5 are valid)



0	1	...	12	13	...	2048	...	2053	...	4608	...
			3000	4				9			

- $PM[2*6] = PM[12] = 3000$ $PM[2*6+1] = PM[13] = 4$
- $PM[4*512+5] = PM[2053] = 9$

Executing VA Translations

- After PM initialization, system processes an input file (separate from the initialization file).
 - The input file consists of a sequence of commands. We will discuss the format of the input file and the commands later.
 - Some of the commands require translating a VA to a PA and writing the result to the output file
 - The result of each translation is either a PA or -1 (error)
- **Translation**
 - break VA into s, p, w, pw
 - if $pw \geq PM[2s]$, then report error; VA is outside of the segment boundary
 - else $PA = PM[PM[2s + 1] * 512 + p] * 512 + w$

Deriving s, p, w, and pw from VA

- s: right-shift VA by 18 bits
 - discards p and w
- w: AND bitwise VA with 9-bit binary constant “1 1111 1111” (1FF)
 - removes all bits other than last 9 — value of w
- p: first right-shift VA by 9 bits to discard w
then AND result with binary constant “1 1111 1111”
 - removes all bits other than the last 9 — value of p
- pw: AND VA with the 18-bit binary constant “11 1111 1111 1111 1111” (3FFFF)
 - removes the leading s component
- Ex: VA = 789002 = 000000011 000000101 000001010, s = 3, p = 5, w = 10

Executing VA Translations - Example

0	1	...	12	13	...	2048	...	2053	...	4608	...
			3000	4				9			

if $pw \geq PM[2s]$, then error

else $PA = PM[PM[2s + 1] * 512 + p] * 512 + w$

$VA = 1575424 = (000000110,000000101,000000000)$

- $s = 6, p = 5, w = 0, pw = 2560, pw < 3000$
- $PA = PM[PM[2*6+1]*512+5]*512+0 = PM[2048+5]*512+0 = 9*512+0 = 4608$

$VA = 1575863, s = 6, p = 5, w = 439, pw = 2999, pw < 3000$

- $PA = \dots = PM[2048+5]*512+439 = 9*512+439 = 5047$

$VA = 1575864, s = 6, p = 5, w = 440, pw = 3000$

- $pw \geq 3000$: error, PA would be 5048 but this is outside of segment boundary

The background of the slide features a collection of open paint cans in various shades of teal, green, and yellow, arranged on a light-colored, textured surface. A paintbrush with a wooden handle and a metal ferrule is positioned horizontally across the middle of the image, its bristles resting on the surface. The overall aesthetic is artistic and painterly.

Extended Version – With Demand Paging

VM with Demand Paging

- Not all pages or PTs are resident in PM
 - must be copied from a paging disk when accessed
 - to avoid page replacement, assume that a free frame is always available
- **Paging Disk**
 - emulated as a two-dimensional integer array, $D[B][512]$
 - B: number of blocks (e.g., 1024)
 - 512: block size (= page size)
- Disk may only be accessed one block at a time:
 - *read_block(b, m)* copies block $D[b]$ into PM frame starting at location $PM[m]$

Extensions

Contents of ST and PT

- If PT is not resident, corresponding ST entry contains a negative number $-b$
 - the absolute value $|-b| = b$ is the block number on disk that contains the PT
- If a page is not resident, corresponding PT entry contains a negative number $-b$
 - b is the block number on disk that contains the page
- The sign bit is used as the present bit (negative = not resident)

Extensions

Free Frames

- Blocks are moved to PM from disk at page fault
- Memory manager must keep track of which frames are free
 - Use a linked list (or mark frames as free)
- Pre-initialization: all frames except frames 0 and 1 are designated as free
 - Frames 0 and 1 are dedicated to segment table, hence not free
- During initialization, frames that are used are no longer free:
 - Frames used to store page tables
 - Frames used to store pages within a segment

VA Translation

- If $pw \geq PM[2s]$, report error; VA is outside of the segment boundary
- If $PM[2s + 1] < 0$, then /* page fault: PT is not resident */
 - Allocate free frame $f1$ using list of free frames
 - Update list of free frames
 - Read disk block $b = |PM[2s + 1]|$ into PM starting at location $f1 * 512$
 - $PM[2s + 1] = f1$ /* update ST entry */
- If $PM[PM[2s + 1] * 512 + p] < 0$ /* page fault: page is not resident */
 - Allocate free frame $f2$ using list of free frames
 - Update list of free frames
 - Read disk block $b = |PM[PM[2s + 1] * 512 + p]|$ into PM starting at $f2 * 512$
 - $PM[PM[2s + 1] * 512 + p] = f2$ /* update PT entry */
- Return $PA = PM[PM[2s + 1] * 512 + p] * 512 + w$

Initialization of PM

Example:

- Line 1: 8 4000 3 9 5000 -7 ...
- Line 2: 8 0 10 8 1 -20 9 0 13 9 1 -25 ...


D

	0	1	...	511	
0					
⋮					
7	13	-25			PT of segment 9
⋮					
20					page 1 of segment 8
⋮					
25					page 1 of segment 9
⋮					

	0	16	17	18	19	1536	1537	5120	6656			
PM		...	4000	3	5000	-7	...	10	-20
	ST (frames 0 and 1)						PT of segment 8 (frame 3)		page 0 of segment 8 (frame 10)		page 0 of segment 9 (frame 13)	


VA Translation: PT and page resident

- $VA = 2097162 = 1000\ 0000000000\ 000001010 = (8, 0, 10)$
- $PM[2*8+1] = 3 \quad (> 0)$: PT is resident at $3*512 = 1536$
- $PM[1536+0] = 10 \quad (> 0)$: pg is resident at $10*512 = 5120$
- $PA = 10*512+10 = 5120+10 = 5130$



VA Translation: PT resident, pg not resident


- $VA = 2097674 = 1000\ 0000000001\ 000001010 = (8, 1, 10)$
- $PM[2*8+1] = 3 \quad (> 0)$: PT is resident at $3*512 = 1536$
- $PM[1536+1] = -20 \quad (< 0)$: pg is in disk block 20
- assume frame 4 is allocated: replace -20 with 4, page 1 starts at $4*512=2048$
- $PA = 2048+10 = 2058$



VA Translation: PT not resident, pg resident

- $VA = 2359306 = 1001\ 0000000000\ 000001010 = (9, 0, 10)$
- $PM[2*9+1] = -7$ (< 0): PT is in disk block 7
- assume frame 5 is allocated: replace -7 with 5
- copy PT from block 7 to $PM[5*512] = 2560$
- $PM[2560+0] = 13$ (> 0): pg is resident at $13*512=6656$
- $PA = 6656+10 = 6666$


	0	1	...	511
0				
:				
7	13	-25		
:				



VA Translation: PT and pg not resident

- $VA = 2359818 = 1001\ 0000000001\ 000001010 = (9, 1, 10)$
- $PM[2*9+1] = -7$ (< 0): PT is in disk block 7
- copy PT from block 7 to $PM[5*512] = 2560$ as before
- $PM[2560+1] = -25$ (< 0): pg is in block 25
- assume frame 14 is allocated, replace -25 by 14
- $PA = 14*512+10 = 7168+10 = 7178$

	0	1	...	511
0				
:				
7	13	-25		
:				



Input file commands

There are three commands that need to be implemented: TA, RP, NL

1. TA (Translate Address):

- TA <va>
 - <va> is an integer, representing a virtual address
 - Translate <va> to a PA and write the result (as a decimal integer) to the output file

2. RP (Read Physical address):

- RP <pa>
 - <pa> is an integer, representing a physical address
 - Read the contents of the physical address <pa> and write the result (as a decimal integer) to the output file

(Continued on next slide)

Input file commands (continued)

3. NL (New Line):

- NL
- Start a new line in the output file.

Output file format:

- A sequence of integers (one for each TA or RP command).
- The integers are to be written one after another on the same line separated by a space,
- A NL command causes a new line to be started.

Example of init, input, output file

init

```
6 3000 4  
6 5 9
```

input

```
TA 1575424  
TA 1575863  
TA 1575864  
NL  
RP 12  
RP 13  
RP 2053  
RP 1024
```

output

```
4608 5047 -1  
3000 4 9 0
```

Summary of Specific Tasks

- Design and implement a VM manager using segmentation and paging
 - Without demand paging (reduced credit), or
 - With demand paging (full credit)
- Skip the TLB
- Design and implement a driver program that
 - Initializes the system from a given initialization file
 - Reads the input file and executes the commands
 - The results are written into an output file